


Master's Thesis in Informatics

Reinforcement Learning for Autonomous Locomotion Control of Snake-Like Robots


Author: Christian Lemke
Christian.Lemke@campus.lmu.de
Supervisor: Prof. Dr.-Ing. habil. Alonis Knoll
Advisor: M.Eng. Zhenshan Bing
Date: 27.07.2018



Content



- Background of Snake-Like Robots
- Background of Reinforcement Learning
- Simulation environment
- Two control experiments
- Conclusion



Snake-Like Robots




Mobility:


- Swim in water
- Climb stairs and poles
- Move in narrow spaces

Use cases:

- Fire fighting
- Inspection and maintenance
- Search and rescue



ACM III (1972)



ACM-R5 (2005)

Snake-Like Robots



The modeling and control problems:

- Highly redundant degrees of freedom
- Complex interaction with the environment
- Difficult to control in real-world situations

Question:

How to solve this complex control problem?

Approach:


Reinforcement Learning

Experiments:

- Autonomous Locomotion Control
- Autonomous Target Tracking



Search and rescue scenario



Box-and-Banana Problem


Proximal Policy Optimization

- Best performance on continuous control tasks
- Easy to implement
- Good sample efficiency
- Small hyperparameter tuning

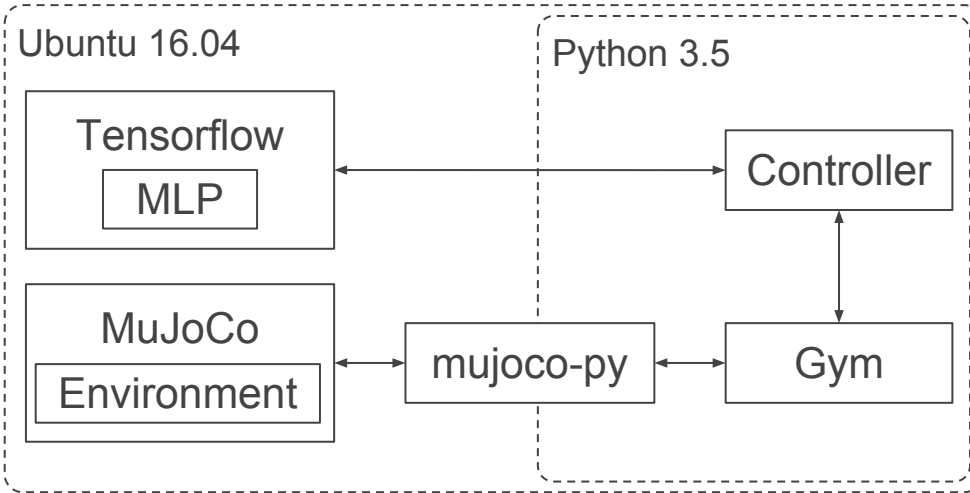
Environment and Robot




- 9 Modules and 8 Joints
- Servo position motors
- Wheels
- Vision via head camera




The robot in MuJoCo



Components overview



Vision of head camera



Rendering with
32x20 RGB

1. Experiment: Autonomous Locomotion Control



Task:


Perform a power efficient locomotion at a specified velocity.

Learn:

- Joint position commands for slithering locomotion
- Control velocity
- Power efficiency

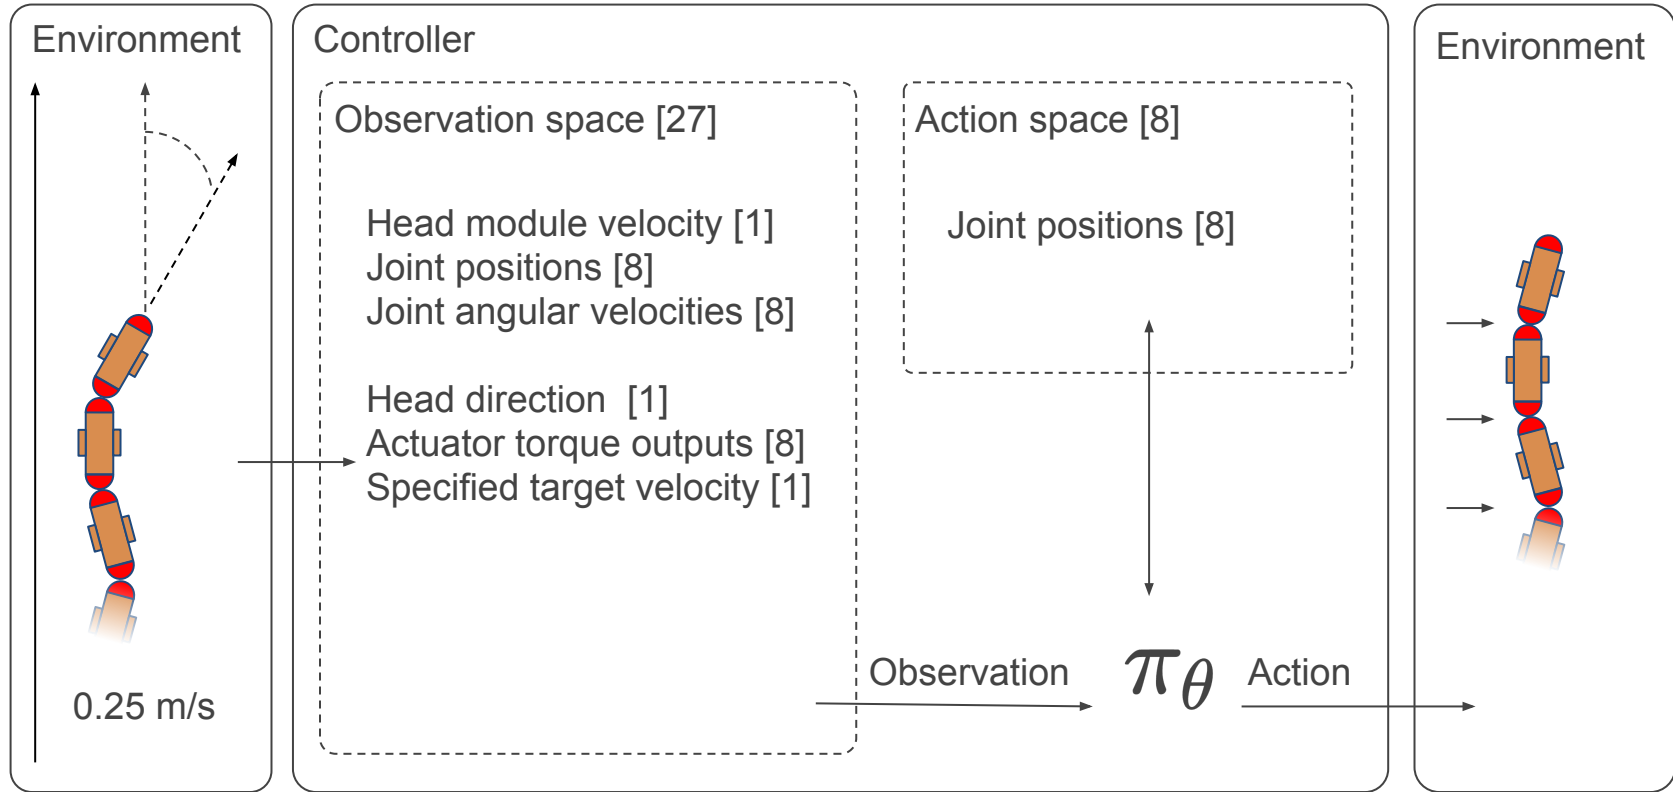
Evaluation:

Comparison to traditional equation controller



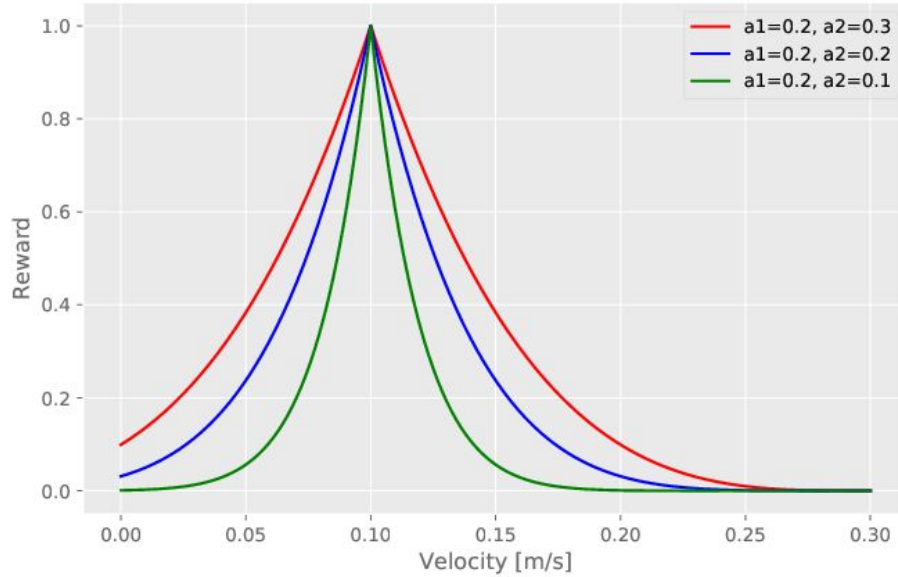
1. Experiment: Autonomous Locomotion Control

Observation Space and Action Space



1. Experiment: Autonomous Locomotion Control

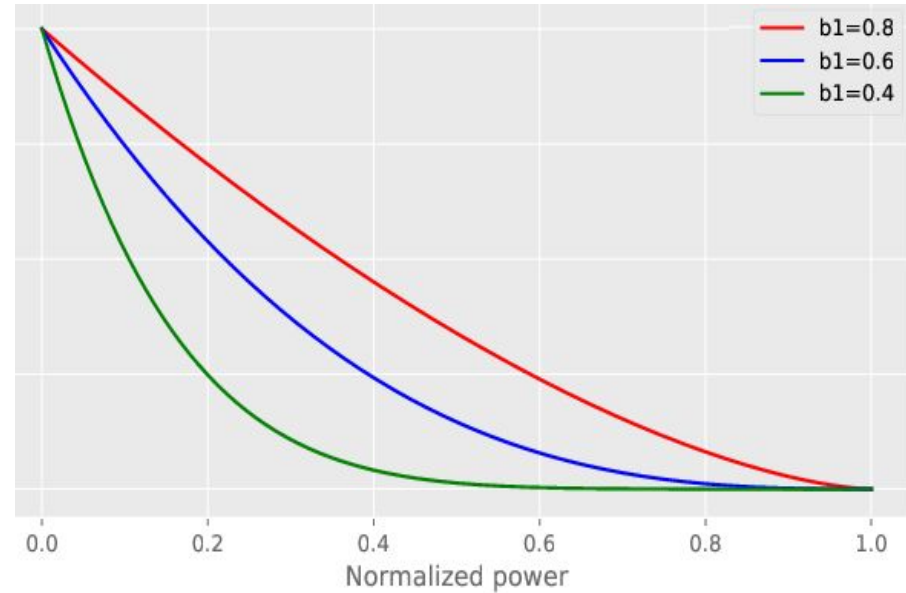
Reward Function



Reward velocity:

$$r_v = \left(1 - \frac{|v_t - v|}{a_1}\right)^{\frac{1}{a_2}}$$

- Reward is combination of velocity and power usage $\mathbf{r} = r_v r_P$




Reward power efficiency:

$$r_P = |1 - \hat{P}|^{b_1^{-2}}$$


$$\hat{P} = \frac{|\tau \dot{\phi}|}{\tau_{max} \dot{\phi}_{max}}$$

1. Experiment: Autonomous Locomotion Control

Results



- Velocity of 0.05 m/s
- “Concertina” gait pattern
- Contracts and stretches the body



- Velocity of 0.25 m/s
- “Lateral undulation” gait pattern
- Carries waves from head to tail

1. Experiment: Autonomous Locomotion Control

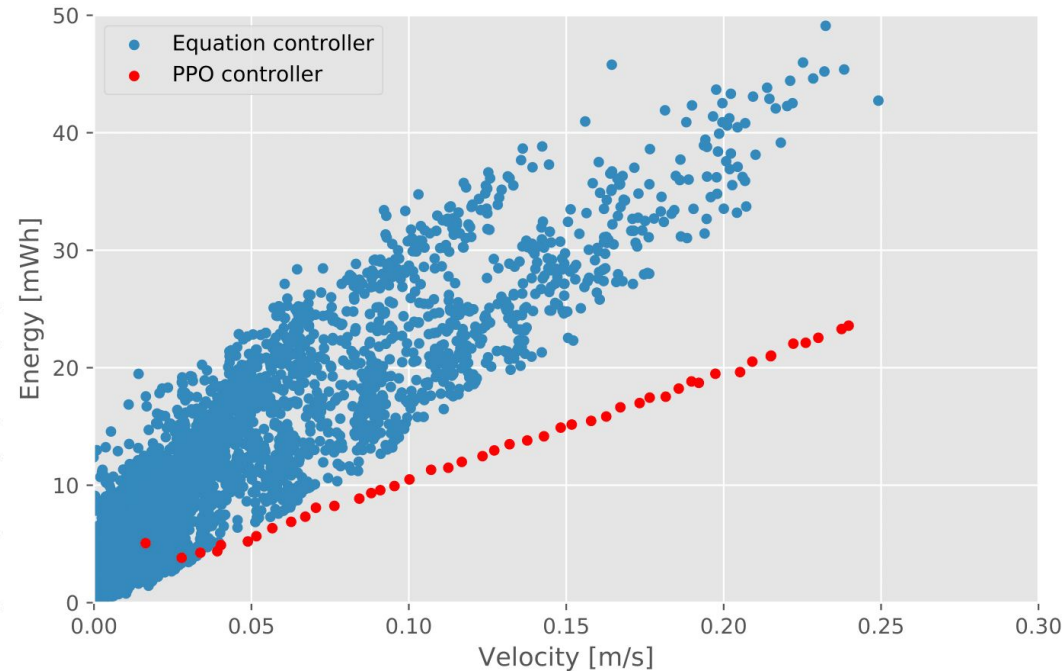
Comparison with traditional Equation Controller



- An efficiency comparison
- Grid search creates a variety of different gaits
- Total of 6480 gait parameter sets

Descriptions	Values
Angular frequency	0.25, 0.5, 0.75, 1.0, 1.25, 1.5, 1.75, 2.0, 2.25, 2.5, 2.75, 3.0
Linear reduction	0.1, 0.2, 0.3, 0.4
Amplitude (in degrees)	40, 50, 60, 70, 80, 90, 100, 110, 120, 130, 140, 150, 160, 170, 180
Bending radius (in degrees)	40, 50, 60, 70, 80, 90, 100, 110, 120

Table of the equation controller parameters



2. Experiment: Autonomous Target Tracking



Task:


Follow a moving target with a certain distance.

Learn:

- Control joints for locomotion
- Use head camera to track the target and estimate the distance



Evaluation:

Test on different target tracks



2. Experiment: Autonomous Target Tracking

Observation Space and Action Space




Select one row:
32x20 to 32 Pixels

2. Experiment: Autonomous Target Tracking







Reward Function



$$r = \frac{|t_d - d_{before}| - |t_d - d_{after}|}{d_r}$$



2. Experiment: Autonomous Target Tracking Result



Conclusion



Advantages:


- Directly solves the problem
- No control engineering

Disadvantages:

- Challenging to develop suitable reward functions
- The policy is difficult to interpret

Future work:

- PPO on a real robot
- 3D Snake-Like Robot Model
- Explore gait adaptiveness on 3D model





Thank you



Master's Thesis in Informatics

Reinforcement Learning for Autonomous Locomotion Control of Snake-Like Robots


Author: Christian Lemke
Christian.Lemke@campus.lmu.de

Backup

Model and Vision




Rendering with
32x20 RGB



Proximal Policy Optimization

- Based on Policy gradient methods
- Best performance on continuous control tasks
- Simple to implement and handle
- Good sample efficiency



A simulated 'humanoid' walker

Proximal Policy Optimization



$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)]$$

Components

Objective function $L^{CLIP}(\theta)$


Probability ratio $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$

Advantage function \hat{A}_t

Min function $\min(value_1, value_2)$

Clip function $\text{clip}(value, min, max)$

Clipping parameter $1 - \epsilon, 1 + \epsilon$



$$L^{CLIP}(\theta) = \hat{E}_t [\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon)\hat{A}_t)]$$

- θ is the policy parameter
- \hat{E}_t denotes the empirical expectation over timesteps
- r_t is the ratio of the probability under the new and old policies, respectively
- \hat{A}_t is the estimated advantage at time t
- ε is a hyperparameter, usually 0.1 or 0.2

1. Experiment: Autonomous Locomotion Control



Reward Function



- Reward is combination of velocity and Power usage
- Both are normalized and multiplied together

$$r_v = \left(1 - \frac{|v_t - v|}{a_1}\right)^{\frac{1}{a_2}}$$

$$r_P = r_{max} |1 - \hat{P}|^{b_1^{-2}}$$



Power measurement



The most straightforward measure of power usage for the j -th actuator ϵ_j is the absolute value of the product of the torque τ_j and its angular velocity $\dot{\phi}_j$. The total power consumption P of all m actuators on each time step is calculated by

$$P = \sum_{j=1}^m |\tau_j \dot{\phi}_j| \quad (4.1)$$

where

$$\tau_j = f_j g_j \quad (4.2)$$

$$\hat{P} = \frac{1}{m} \sum_{j=1}^m \frac{|f_j g_j \dot{\phi}_j|}{f_{max} g_j \dot{\phi}_{max}}$$

1. Experiment: Autonomous Locomotion Control

Equation controller



Traditionally locomotion gait




Table 5.1: Grid search parameters for the equation controller

Parameters	Descriptions	Values
w	Angular frequency	0.25, 0.5, 0.75, 1.0, 1.25, 1.5, 1.75, 2.0, 2.25, 2.5, 2.75, 3.0
y (x=1-y)	Linear reduction	0.1, 0.2, 0.3, 0.4
alpha	Amplitude (in degrees)	40, 50, 60, 70, 80, 90, 100, 110, 120, 130, 140, 150, 160, 170, 180
lambda	Bending radius (in degrees)	40, 50, 60, 70, 80, 90, 100, 110, 120

Table 5.3: Overview of the PPO controller observation space parameters



Symbols	Descriptions
ϕ_{1-8}	Relative joint angular positions
$\dot{\phi}_{1-8}$	Relative joint angular velocity
v_1	Absolute head module velocity (measured at (x_1, y_1))
τ_{1-8}	Actuator torque output
ϕ_t	Relative angle between the head direction and the target
v_t	Specified target velocity

Table 6.2: Overview of the target tracking controller observation space parameters

Symbols	Descriptions
ϕ_{1-8}	Relative joint angular positions
$\dot{\phi}_{1-8}$	Relative joint angular velocity
v_1	Head link velocity (measured at (x_1, y_1))
$p_{10,1-32}$	Pixel 1 to 32 of the 10th row of the gray camera image



2. Experiment: Autonomous Target Tracking

Result: Zigzag-Track (1/4)



2. Experiment: Autonomous Target Tracking

Result: Line-Track (2/4)






Mean: 3.99 m

Std: 0.18 m

2. Experiment: Autonomous Target Tracking

Result: Wave-Track (3/4)






Mean: 3.99 m

Std: 0.24 m

2. Experiment: Autonomous Target Tracking


Result: Random-Track (4/4)




Mean: 3.88 m

Std: 0.22 m


2. Experiment: Autonomous Target Tracking Result



- Stops to keep distance
- Uses different
- Path following behavior
- Does not face the target directly




Learning process



1. Experiment: Autonomous Locomotion Control


Reward Function



$$r_v = \left(1 - \frac{|v_t - v|}{a_1}\right)^{\frac{1}{a_2}}$$

1. Experiment: Autonomous Locomotion Control

Reward Function



$$r_P = r_{max} |1 - \hat{P}|^{b_1^{-2}}$$

Training 1





Figure 5.5: This plot shows the learning curve of the PPO controller. Mean episode reward over 3 million time steps. The x-axes represent the number of network updates and the y-axis the achieved mean episode reward per update. Note that an update contains 2048 time steps.

Training 2




Figure 6.4: The learning curve of the autonomous target tracking model. It is trained with 3 million time steps with 1000 time steps per episode and 1024 time steps per update. The mean episode reward does not converge to a specific reward.